

COL380

Introduction to
Parallel & Distributed Programming


Agenda

- Task graphs and parallel programming models
- Shared memory programming
- Concurrency and race conditions
- Memory and Consistency

- Multiple interacting fragments
 - ➔ Shared state
 - ➔ Fragments may have private (hidden) state
- Usually Asynchronous
 - ➔ Synchronous models can simplify some things

Parallel Programs


- Multiple interacting fragments
 - ➔ Shared state
 - ➔ Fragments may have private (hidden) state
- Usually Asynchronous
 - ➔ Synchronous models can simplify some things



Program can be represented as a graph

Parallel Programs


- Multiple interacting fragments
 - ➔ Shared state
 - ➔ Fragments may have private (hidden) state
- Usually Asynchronous
 - ➔ Synchronous models can simplify some things



Program can be represented as a graph


Parallel Programs

- Multiple interacting fragments
 - ➔ Shared state
 - ➔ Fragments may have private (hidden) state
- Usually Asynchronous
 - ➔ Synchronous models can simplify some things



Program can be represented as a graph

Task Graph



- Granularity
- Critical Path Length
- Maximum Concurrency
- Average Concurrency

$$= \frac{\text{\#tasks}}{\text{Critical path length}}$$

- Shared Memory model
- Distributed Memory/Message passing model
- Task-graph based model
- Fork-join model
- Work-queue model
- Stream processing model
- Map-reduce model
- Client-server model

- Shared Memory model
- Distributed Memory/Message passing model
- Task-graph based model
- Fork-join model
- Work-queue model
- Stream processing model
- Map-reduce model
- Client-server model

```
void processAB (int a[], int b[], int n)
{
    for(int i=0; i<n; i += 3) {
        #pragma omp task
        subtaskA(a+i);
        #pragma omp task
        subtaskB(b+i);
    }
}

.....

#pragma omp task
processAB(a, b, n);
```

- Shared Memory model
- Distributed Memory/Message passing model
- Task-graph based model
- Fork-join model
- Work-queue model
- Stream processing model
- Map-reduce model
- Client-server model

```
void processAB (int a[], int b[], int n)
{
    for(int i=0; i<n; i += 3) {
        #pragma omp task depend(out: x)
        subtaskA(a+i);
        #pragma omp task depend(in: x)
        subtaskB(b+i);
    }
}

.....

#pragma omp task
processAB(a, b, n);
```

- **Shared Memory**

- ➔ Tasks share a common address space they access **asynchronously**
 - ▶ Can **Synchronize** to enforce certain order
- ➔ Implicit interaction through memory
- * Data may be cached on the processor executing the task

- **Message Passing**

- ➔ Tasks have their own address spaces
- ➔ Explicit Inter-process interaction (in code): <send> and <receive>


```
void loop(int &a, int n)
{
    for(int i=0; i<n; i += 3)
        a++;
}
```


(Rd a; Add1; Wr a)

```
..
int a = 0, n = 10;
#pragma omp parallel
    loop(a, n);
```


Sharing and Racing

```
void loop(int &a, int n)
{
    for(int i=0; i<n; i += 3)
        a++;
}


..
int a = 0, n = 10;
#pragma omp parallel
    loop(a, n);
```

- Operation are **not instantaneous**: 
- Such an operation can becomes visible to different threads at different times
- apparent order of operations may not be *consistent*
- Even *consistent* operations can be *concurrent* (vary from execution to execution)
- Program must remain correct no matter the order

Sharing and Racing

```
void loop(int &a, int n)
{
    for(int i=0; i<n; i += 3)
        a++;
}
```

```
..
int a = 0, n = 10;
#pragma omp parallel
    loop(a, n);
```

- Operation are **not instantaneous**: 
- Such an operation can becomes visible to different threads at different times
- apparent order of operations may not be *consistent*
- Even *consistent* operations can be *concurrent* (vary from execution to execution)
- Program must remain correct no matter the order

Data race: Concurrent RW or WW operations

Race condition: If variable order of concurrent operations affects correctness

```
void loop(int &a, int n)
{
    for(int i=0; i<n; i += 3)
        a++;
}
```

(Rd a; Add1; Wr a)

```
..
int a = 0, n = 10;
#pragma omp parallel
    loop(a, n);
```

```
    movl    (%rdi), %eax
    addl    $1, %eax
    movl    $0, %edx
.L3: movl    %eax, %ecx
    addl    $3, %edx
    addl    $1, %eax
    cmpl    %edx, %esi
    jg      .L3
    movl    %ecx, (%rdi)
```



```
void loop(volatile int &a, int n)
{
    for(int i=0; i<n; i += 3)
        a++;
}
```

(Rd a; Add1; Wr a)

```
..
int a = 0, n = 10;
#pragma omp parallel
    loop(a, n);
```

```
    movl    (%rdi), %eax
    addl    $1, %eax
    movl    $0, %edx
.L3: movl    %eax, %ecx
    addl    $3, %edx
    addl    $1, %eax
    cmpl    %edx, %esi
    jg      .L3
    movl    %ecx, (%rdi)
```



```
void loop(volatile int &a, int n)
{
    for(int i=0; i<n; i += 3)
        a++;
}
(Rd a; Add1; Wr a)

..
int a = 0, n = 10;
#pragma omp parallel
    loop(a, n);
```

```
    movl    (%rdi), %eax
    addl    $1, %eax
    movl    $0, %edx
.L3: movl    %eax, %ecx
    addl    $3, %edx
    addl    $1, %eax
    cmpl    %edx, %esi
    jg      .L3
    movl    %ecx, (%rdi)
```

```
.L3: movl    (%rdi), %eax
    addl    $1, %eax
    movl    %eax, (%rdi)
    addl    $3, %edx
    cmpl    %edx, %esi
    jg      .L3
    ret
```

```
void loop(volatile int &a, int n)
{
    for(int i=0; i<n; i += 3)
        a++;
}
```

(Rd a; Add1; Wr a)

```
..
int a = 0, n = 10;
#pragma omp parallel
    loop(a, n);
```

```
    movl    (%rdi), %eax
    addl    $1, %eax
    movl    $0, %edx
.L3: movl    %eax, %ecx
    addl    $3, %edx
    addl    $1, %eax
    cmpl    %edx, %esi
    jg      .L3
    movl    %ecx, (%rdi)
```

```
.L3: movl    (%rdi), %eax
    addl    $1, %eax
    movl    %eax, (%rdi)
    addl    $3, %edx
    cmpl    %edx, %esi
    jg      .L3
    ret
```

var = val need not be “seen”
or, need not be “atomically” seen

- Task graphs and parallel programming models
- Shared memory programming
- Concurrency and race conditions
- Memory and Consistency